
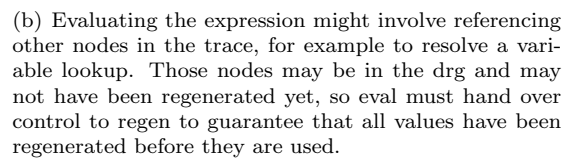
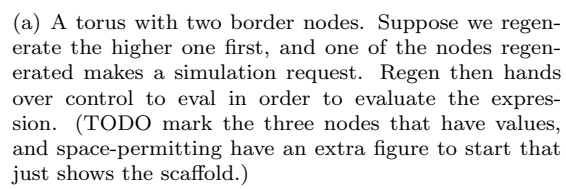


(a) First we walk downstream from the principal node, and color gold every node whose value may change, and blue every node at which we can absorb.




(b) Next, we color green every node that may no longer exist once the gold nodes are resampled. At this point, the red and gold nodes constitute the definite regeneration graph, the blue nodes constitute the absorbing border, and the green nodes constitute the brush.


Figure 1: The two stages of constructing a scaffold



2




(a) A large scaffold for sampling a hyperparameter.




(b) The application of the maker SP computes the log density of all of its applications for us. We say that the maker SP “absorbs at applications”.

Figure 3: Absorbing at applications (AAA)



(a) A scaffold with three border nodes, one of which is absorbing at applications.



(b) A simulation request may lookup a node that is itself a reference to the aaa node. Even though regen will be called on a node that is not in the drg, the aaa node it refers to must be regenerated nonetheless.

Figure 4: Challenges with absorbing at applications